

Technology in Medicine

By: Lauren Esser



BACKGROUND

Pneumonia is a lung infection that can cause mild to severe illness in people of all ages.

Currently to diagnose pneumonia the doctor will ask about your medical history and may perform a blood test, chest x-ray, pulse oximetry, or sputum test.



FACTS ABOUT PNEUMONIA



VACCINE

- There is a vaccine to help fight against pneumonia, 68.9% of adults over the age of 65 have received the vaccine.



HOSPITAL CARE

- More than 250,000 people have to seek hospital care each year due to pneumonia.



TREATMENTS

- Current treatments are antibiotics, cough medicine, fever reducers, pain relievers, and potentially hospital care.




DEATHS

- Pneumonia kills more children younger than 5 years old each year than any other infectious disease.
- About 50,000 people die from the disease each year in the U.S.


PROBLEM STATEMENT

Build a neural network that classifies x-ray images of pediatric patients to identify whether or not they have pneumonia.

Normal X-Ray



Pneumonia X-Ray





DATASET

From Kaggle.com titled
Chest X-Ray Images
(Pneumonia)

	TRAIN	TEST	VALIDATION
NORMAL	1341	234	8
PNEUMONIA	3875	390	8

OSEMN



OBTAIN DATA

Download from Kaggle. Import to Google Drive and unzip file for access.



SCRUB DATA

Organize data into proper directories and check length.



EXPLORE DATA

Visualizations, prepare model by reshaping and resizing images.



MODEL DATA

First create baseline Neural Network models and work up to CNN Models.



INTERPRET DATA

Explain best model use.

METHODS USED



CNN



Load Images & Design Model



Alternate convolutional and pooling layers.



Include more parameters in later layers.



Add final dense later to add a classifier to the convolutional base

**METHODS
USED**



My Recommendation: CNN Model 3

Model: "sequential_16"

Layer (type)	Output Shape	Param #
conv2d_44 (Conv2D)	(None, 62, 62, 64)	1792
max_pooling2d_44 (MaxPooling)	(None, 31, 31, 64)	0
conv2d_45 (Conv2D)	(None, 29, 29, 128)	73856
max_pooling2d_45 (MaxPooling)	(None, 14, 14, 128)	0
conv2d_46 (Conv2D)	(None, 12, 12, 512)	590336
max_pooling2d_46 (MaxPooling)	(None, 6, 6, 512)	0
conv2d_47 (Conv2D)	(None, 4, 4, 512)	2359808
max_pooling2d_47 (MaxPooling)	(None, 2, 2, 512)	0
flatten_12 (Flatten)	(None, 2048)	0
dense_61 (Dense)	(None, 128)	262272
dense_62 (Dense)	(None, 64)	8256
dense_63 (Dense)	(None, 32)	2080
dense_64 (Dense)	(None, 1)	33
Total params: 3,298,433		
Trainable params: 3,298,433		
Non-trainable params: 0		

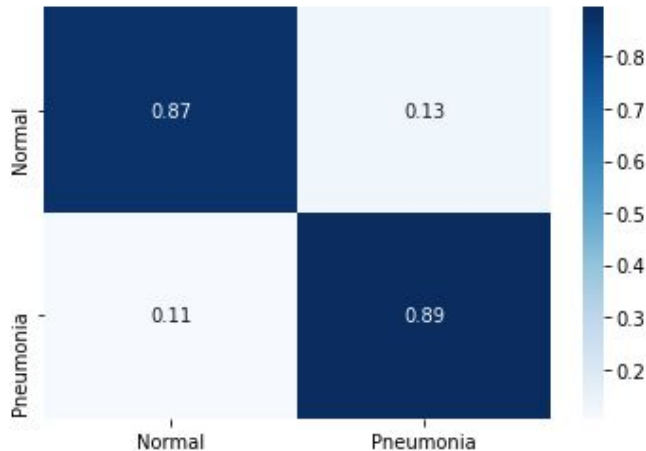
Reasoning:

1. High Accuracy
2. High Recall
3. High Precision

CLASSIFICATION REPORT

	precision	recall	f1-score	support
Normal	0.83	0.87	0.85	234
Pneumonia	0.92	0.89	0.91	390
accuracy			0.88	624
macro avg	0.88	0.88	0.88	624
weighted avg	0.89	0.88	0.89	624

CONFUSION MATRIX



1 Use a Convolutional Neural Network over a Basic Neural Network.

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 64)	786496
dense_1 (Dense)	(None, 32)	2080
dense_2 (Dense)	(None, 1)	33

Total params: 788,609

Trainable params: 788,609


Non-trainable params: 0

Basic Neural Network Model

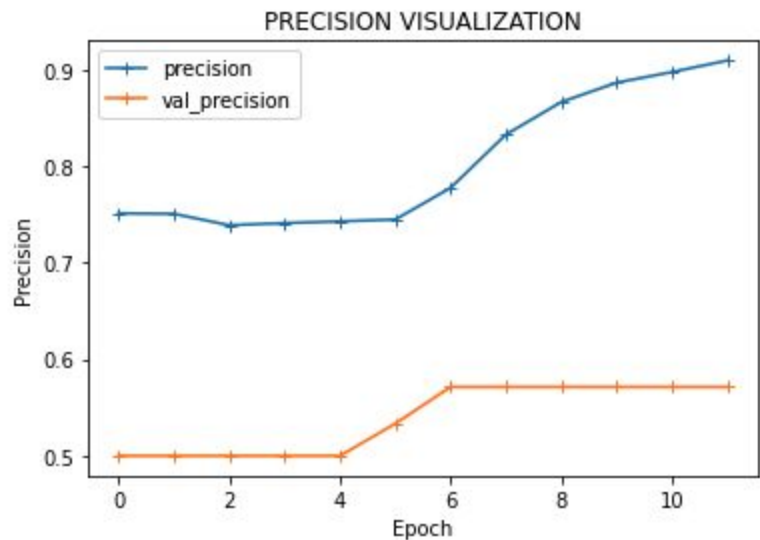
CLASSIFICATION REPORT

	precision	recall	f1-score	support
Normal	0.78	0.74	0.76	234
Pneumonia	0.85	0.87	0.86	390
accuracy			0.82	624
macro avg	0.81	0.81	0.81	624
weighted avg	0.82	0.82	0.82	624

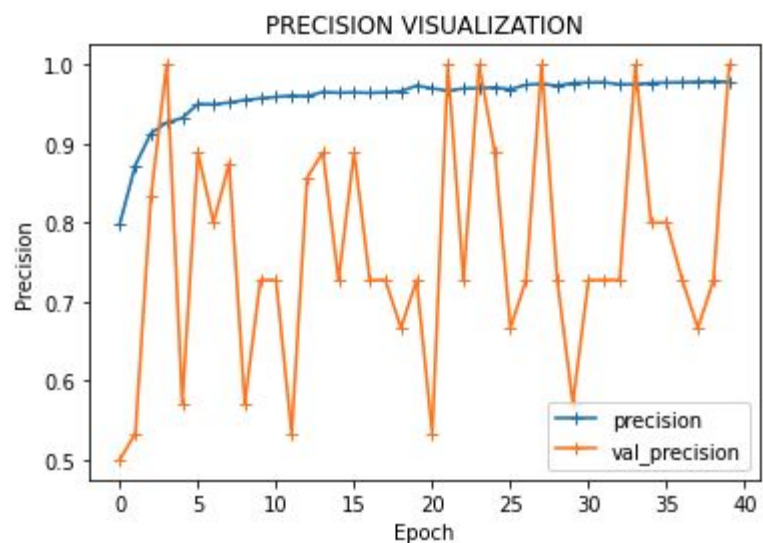
CONFUSION MATRIX



The Basic Neural Network has lower accuracy, precision, and recall



Convolutional
Neural Network



Basic Neural
Network

2 IMAGE SIZE

64 x 64 image size took 19.6 seconds to run a baseline neural network model. 82% accuracy

150 x 150 image size took 41.3 seconds to run the same baseline neural network model. 80% accuracy

See Appendix One A & B for Images of runtime/accuracy

3 Precision or Recall?

- CNN Model 2 had a Pneumonia recall of 100%
- Recall states how many of the actual cases the model found.

```
-----  
CLASSIFICATION REPORT  
              precision    recall  f1-score   support  
  
   Normal      0.99      0.33      0.49      234  
  Pneumonia    0.71      1.00      0.83      390  
  
 accuracy      0.75      0.75      0.75      624  
 macro avg     0.85      0.66      0.66      624  
 weighted avg  0.82      0.75      0.70      624  
-----
```

Recall / Sensitivity

$$\frac{TP}{(TP + FN)}$$

3 Precision or Recall?

- CNN Model 3 the highest precision at 92%
- Precision states the percentage of positive predictions the model stated correctly.

CLASSIFICATION REPORT

	precision	recall	f1-score	support
Normal	0.83	0.87	0.85	234
Pneumonia	0.92	0.89	0.91	390
accuracy			0.88	624
macro avg	0.88	0.88	0.88	624
weighted avg	0.89	0.88	0.89	624

Precision

$$\frac{TP}{(TP + FP)}$$

3 Precision or Recall: Depends on the Case

Choose Precision if...

The cure is a tough or dangerous processes. Example: You don't want to give someone chemo or radiation when they don't actually have the illness.

Choose Recall if...

Miss identifying leads to more deaths. Recall will help identify more people as sick.


1 Use a Convolutional Neural Network over a Basic Neural Network.

2 Use a 64 x 64 image size when using the model.

3 Precision vs. Recall: Depending on the illness/treatment one may be more important than the other. In the case of pneumonia focus on recall.

SUMMARY OF
RECOMMENDATIONS





Future Recommendations



- 1 Use a larger dataset to improve accuracy and precision
- 2 Look at X-Rays of adult lungs with and without pneumonia
- 3 Use premade CNN to see if they could potentially work better.

THANK YOU FOR YOUR TIME!

Questions?

Feel free to contact me at:
Lauren.Esser02@gmail.com
or
Lauren Esser on LinkedIn

APPENDIX 1A: 64x64

```
Epoch 37/40
105/105 [=====] - 0s 4ms/step - loss: 0.0906 - acc: 0.9655
Epoch 38/40
105/105 [=====] - 0s 4ms/step - loss: 0.0925 - acc: 0.9670
Epoch 39/40
105/105 [=====] - 0s 4ms/step - loss: 0.0885 - acc: 0.9691
Epoch 40/40
105/105 [=====] - 0s 4ms/step - loss: 0.0854 - acc: 0.9674
Time to Run: 0:00:21.512352
```

----- CLASSIFICATION REPORT

	precision	recall	f1-score	support
Normal	0.76	0.75	0.76	234
Pneumonia	0.85	0.86	0.85	390
accuracy			0.82	624
macro avg	0.81	0.80	0.80	624
weighted avg	0.82	0.82	0.82	624

APPENDIX 1B: 150x150

```
Epoch 37/40
105/105 [=====] - 1s 10ms/step - loss: 0.1022 - acc: 0.9618
Epoch 38/40
105/105 [=====] - 1s 9ms/step - loss: 0.1081 - acc: 0.9584 -
Epoch 39/40
105/105 [=====] - 1s 10ms/step - loss: 0.1000 - acc: 0.9649
Epoch 40/40
105/105 [=====] - 1s 10ms/step - loss: 0.0981 - acc: 0.9653
Time to Run: 0:00:41.168897
```

CLASSIFICATION REPORT

	precision	recall	f1-score	support
Normal	0.81	0.63	0.71	234
Pneumonia	0.80	0.91	0.85	390
accuracy			0.80	624
macro avg	0.81	0.77	0.78	624
weighted avg	0.80	0.80	0.80	624



APPENDIX 2

Model: "sequential_16"

Layer (type)	Output Shape	Param #
conv2d_44 (Conv2D)	(None, 62, 62, 64)	1792
max_pooling2d_44 (MaxPooling)	(None, 31, 31, 64)	0
conv2d_45 (Conv2D)	(None, 29, 29, 128)	73856
max_pooling2d_45 (MaxPooling)	(None, 14, 14, 128)	0
conv2d_46 (Conv2D)	(None, 12, 12, 512)	590336
max_pooling2d_46 (MaxPooling)	(None, 6, 6, 512)	0
conv2d_47 (Conv2D)	(None, 4, 4, 512)	2359808
max_pooling2d_47 (MaxPooling)	(None, 2, 2, 512)	0
flatten_12 (Flatten)	(None, 2048)	0
dense_61 (Dense)	(None, 128)	262272
dense_62 (Dense)	(None, 64)	8256
dense_63 (Dense)	(None, 32)	2080
dense_64 (Dense)	(None, 1)	33

Total params: 3,298,433

Trainable params: 3,298,433

Non-trainable params: 0



APPENDIX 3

Model: "sequential_5"

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 62, 62, 64)	1792
max_pooling2d_2 (MaxPooling2D)	(None, 31, 31, 64)	0
conv2d_3 (Conv2D)	(None, 29, 29, 96)	55392
max_pooling2d_3 (MaxPooling2D)	(None, 14, 14, 96)	0
flatten_1 (Flatten)	(None, 18816)	0
dense_17 (Dense)	(None, 64)	1204288
dense_18 (Dense)	(None, 32)	2080
dense_19 (Dense)	(None, 16)	528
dense_20 (Dense)	(None, 1)	17

Total params: 1,264,097

Trainable params: 1,264,097

Non-trainable params: 0

APPENDIX 4

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 64)	786496
dense_1 (Dense)	(None, 32)	2080
dense_2 (Dense)	(None, 1)	33

Total params: 788,609

Trainable params: 788,609

Non-trainable params: 0

APPENDIX 5A

Model: "sequential_16"

Layer (type)	Output Shape	Param #
conv2d_44 (Conv2D)	(None, 62, 62, 64)	1792
max_pooling2d_44 (MaxPooling)	(None, 31, 31, 64)	0
conv2d_45 (Conv2D)	(None, 29, 29, 128)	73856
max_pooling2d_45 (MaxPooling)	(None, 14, 14, 128)	0
conv2d_46 (Conv2D)	(None, 12, 12, 512)	590336
max_pooling2d_46 (MaxPooling)	(None, 6, 6, 512)	0
conv2d_47 (Conv2D)	(None, 4, 4, 512)	2359808
max_pooling2d_47 (MaxPooling)	(None, 2, 2, 512)	0
flatten_12 (Flatten)	(None, 2048)	0
dense_61 (Dense)	(None, 128)	262272
dense_62 (Dense)	(None, 64)	8256
dense_63 (Dense)	(None, 32)	2080
dense_64 (Dense)	(None, 1)	33

Total params: 3,298,433

Trainable params: 3,298,433

Non-trainable params: 0

APPENDIX 5B

